

## ▼ FINAL PROJECT

```
from google.colab import files
uploaded = files.upload()
```



Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving test-Project1.csv to test-Project1.csv  
Saving train-Project1.csv to train-Project1.csv

```
#Importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```



/usr/local/lib/python3.6/dist-packages/statsmodels/tools/\_testing.py:19: FutureWarning  
import pandas.util.testing as tm

```
#importing datasets
data_train = pd.read_csv('train-Project1.csv')
data_test = pd.read_csv('test-Project1.csv')
data_train.head()
```



	Id	Open Date	City	City Group	Type	P1	P2	P3	P4	P5	P6	P7	P8	P9	P10	P
0	0	07/17/1999	İstanbul	Big Cities	IL	4	5.0	4.0	4.0	2	2	5	4	5	5	
1	1	02/14/2008	Ankara	Big Cities	FC	4	5.0	4.0	4.0	1	2	5	5	5	5	
2	2	03/09/2013	Diyarbakır	Other	IL	2	4.0	2.0	5.0	2	3	5	5	5	5	
3	3	02/02/2012	Tokat	Other	IL	6	4.5	6.0	6.0	4	4	10	8	10	10	
4	4	05/09/2009	Gaziantep	Other	IL	3	4.0	3.0	4.0	2	2	5	5	5	5	

## ▼ Visualizing the dataset

```
cities = data_train['City'].unique()
cities.sort()

groups = data_train['City Group'].unique()

types = data_train['Type'].unique()
```

```
citywisedata = data_train.groupby('City').mean()
citywisedata.head()
```



	Id	P1	P2	P3	P4	P5	P6
City							
<b>Adana</b>	72.666667	4.000000	5.000000	4.000000	3.000000	1.000000	2.666667
<b>Afyonkarahisar</b>	8.000000	1.000000	1.000000	4.000000	4.000000	1.000000	2.000000
<b>Amasya</b>	110.000000	6.000000	3.000000	6.000000	6.000000	4.000000	4.000000
<b>Ankara</b>	62.105263	3.526316	4.421053	4.078947	4.736842	2.789474	3.947368
<b>Antalya</b>	74.000000	3.000000	3.500000	5.250000	4.375000	2.000000	4.000000

```
#Plotting the mean data against each city
fig, ax = plt.subplots()
fig.set_size_inches(10, 10)
ax.set_xticklabels(ax.get_xticklabels(), rotation=40, ha="right")
sns.barplot(x = cities, y = citywisedata['revenue'], ax = ax)
```



<matplotlib.axes.\_subplots.AxesSubplot at 0x7fd4357c47f0>




```
#Plotting the city type against its mean revenue
citytypewisedata = data_train.groupby('City Group').mean()
```

```
fig, ax = plt.subplots()
fig.set_size_inches(10, 10)
```

```
sns.barplot(x = groups, y = citytypewisedata['revenue'], ax = ax)
```



<matplotlib.axes.\_subplots.AxesSubplot at 0x7fd43560edd8>




```
#Plotting each city group against its mean
citygroupwisedata = data_train.groupby('Type').mean()
```

```
fig, ax = plt.subplots()
fig.set_size_inches(10, 10)
```

```
sns.barplot(x = types, y = citygroupwisedata['revenue'], ax = ax)
```



&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7fd4350d43c8&gt;



```
cat_col=data_train.select_dtypes(include='object').columns  
num_col=data_train.select_dtypes(exclude='object').columns
```

```
#Plotting the categorical data against the target values
```

```
sns.boxplot(x = 'City Group', y = 'revenue', data = data_train)
```




```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd4355c9710>
```

```
sns.boxplot(x = 'Type', y = 'revenue', data = data_train)
```




```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd435039f28>
```




#Plotting the Numerical data against the Target Variable


```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P1', 'P2', 'P3'])
```




```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P4', 'P5', 'P6', 'P7'])
```




```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P8', 'P9', 'P10', 'P11'])
```



```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P12', 'P13', 'P14', 'P15'])
```




```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P16', 'P17', 'P18', 'P19'])
```




```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P20', 'P21', 'P22', 'P23'])
```






```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P24', 'P25', 'P26', 'P27'])
```



```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P28', 'P29', 'P30', 'P31', 'P32'])
```



```
pp = sns.pairplot(data=data_train,
                  y_vars=['revenue'],
                  x_vars=['P33', 'P34', 'P35', 'P36', 'P37'])
```



## ▼ Preprocessing the dataset

```
#Creating a flag for each type of restaurant
data_train['Type_IL'] = np.where(data_train['Type'] == 'IL', 1, 0)
data_train['Type_FC'] = np.where(data_train['Type'] == 'FC', 1, 0)
data_train['Type_DT'] = np.where(data_train['Type'] == 'DT', 1, 0)

#Creating a flag for 'Big Cities'
data_train['Big_Cities'] = np.where(data_train['City Group'] == 'Big Cities', 1, 0)

#Converting Open_Date into day count
#Considering the same date the dataset was made available
data_train['Days_Open'] = (pd.to_datetime('2015-03-23') - pd.to_datetime(data_train['Open Date']).dt.days)

#Removing unused columns
data_train = data_train.drop('Type', axis=1)
data_train = data_train.drop('City Group', axis=1)
data_train = data_train.drop('City', axis=1)
data_train = data_train.drop('Open Date', axis=1)

#Adjusting test data as well
data_test['Type_IL'] = np.where(data_test['Type'] == 'IL', 1, 0)
data_test['Type_FC'] = np.where(data_test['Type'] == 'FC', 1, 0)
data_test['Type_DT'] = np.where(data_test['Type'] == 'DT', 1, 0)
data_test['Big_Cities'] = np.where(data_test['City Group'] == 'Big Cities', 1, 0)
data_test['Days_Open'] = (pd.to_datetime('2015-03-23') - pd.to_datetime(data_test['Open Date']).dt.days)
data_test = data_test.drop('Type', axis=1)
data_test = data_test.drop('City Group', axis=1)
data_test = data_test.drop('City', axis=1)
data_test = data_test.drop('Open Date', axis=1)
data_train.head()
```




	Id	P1	P2	P3	P4	P5	P6	P7	P8	P9	P10	P11	P12	P13	P14	P15	P16	P17	P18
0	0	4	5.0	4.0	4.0	2	2	5	4	5	5	3	5	5.0	1	2	2	2	2
1	1	4	5.0	4.0	4.0	1	2	5	5	5	5	1	5	5.0	0	0	0	0	0
2	2	2	4.0	2.0	5.0	2	3	5	5	5	5	2	5	5.0	0	0	0	0	0
3	3	6	4.5	6.0	6.0	4	4	10	8	10	10	8	10	7.5	6	4	9	3	3
4	4	3	4.0	3.0	4.0	2	2	5	5	5	5	2	5	5.0	2	1	2	1	1

## ▼ Implementing the models



```
X = data_train.drop(['Id', 'revenue'], axis=1)
Y = data_train.revenue
X.shape
```

 (137, 42)

```
#implementing ols regressor and plotting its summary
import statsmodels.api as sm
X_temp = X
X_temp = np.append(arr = np.ones((137,1)).astype(int), values = X, axis = 1)

X_temp = X_temp.astype(np.float64)
regressor_OLS = sm.OLS(Y, X_temp).fit()
regressor_OLS.summary()
```



```

x5  1.294e+05  3.91e+05 0.331  0.742 -6.47e+05 9.06e+05
x6  2.867e+05  2.68e+05 1.071  0.287 -2.45e+05 8.18e+05
x7  -1.127e+04 2.85e+05 -0.040 0.969 -5.77e+05 5.54e+05
x8  -1.203e+06 5.38e+05 -2.236 0.028 -2.27e+06 -1.35e+05
x9  1.488e+06  1.06e+06 1.400  0.165 -6.22e+05 3.6e+06
x10 -6.402e+04 1.67e+06 -0.038 0.970 -3.38e+06 3.25e+06
x11 -3.063e+05 3.07e+05 -0.999 0.320 -9.15e+05 3.02e+05
x12 -3.13e+05  6.88e+05 -0.455 0.650 -1.68e+06 1.05e+06
x13 -5.434e+05 1.57e+06 -0.347 0.729 -3.65e+06 2.57e+06
x14 -1.194e+05 3.52e+05 -0.339 0.736 -8.19e+05 5.8e+05
x15 -2.158e+05 4.94e+05 -0.436 0.664 -1.2e+06  7.66e+05
x16 -4.039e+05 5.58e+05 -0.724 0.471 -1.51e+06 7.04e+05
x17 2.054e+05  3.71e+05 0.553  0.582 -5.32e+05 9.43e+05
x18 3.361e+05  4.15e+05 0.811  0.420 -4.87e+05 1.16e+06
x19 -8.985e+04 1.45e+05 -0.620 0.536 -3.77e+05 1.98e+05
x20 -3.053e+05 1.82e+05 -1.680 0.096 -6.66e+05 5.55e+04
x21 1.497e+05  2.6e+05  0.576  0.566 -3.66e+05 6.65e+05
x22 -2.983e+05 2.72e+05 -1.095 0.276 -8.39e+05 2.42e+05
x23 1.523e+05  1.33e+05 1.149  0.254 -1.11e+05 4.16e+05
x24 5.087e+05  5.69e+05 0.895  0.373 -6.2e+05  1.64e+06
x25 3.712e+05  5.49e+05 0.676  0.500 -7.18e+05 1.46e+06
x26 -1.071e+06 6.13e+05 -1.748 0.084 -2.29e+06 1.46e+05
x27 9.661e+04  2.28e+05 0.423  0.673 -3.57e+05 5.5e+05
x28 4.75e+05   3.2e+05  1.484  0.141 -1.6e+05  1.11e+06
x29 -1.054e+05 3.42e+05 -0.308 0.759 -7.85e+05 5.74e+05
x30 6.346e+04  1.7e+05  0.373  0.710 -2.74e+05 4.01e+05
x31 1.609e+05  2.73e+05 0.590  0.557 -3.81e+05 7.02e+05
x32 -3.3e+05   2.66e+05 -1.241 0.218 -8.58e+05 1.98e+05

```

#Implementing Multiple Regression

```

from sklearn.linear_model import LinearRegression
from sklearn import metrics
regressor = LinearRegression()
regressor.fit(X, Y)

```

```

test_predicted_mreg = pd.DataFrame()
test_predicted_mreg['Id'] = data_test.Id
test_predicted_mreg['Prediction'] = regressor.predict(data_test.drop('Id', axis=1))
test_predicted_mreg.head()

```



	Id	Prediction
0	0	4.669238e+06
1	1	2.741557e+06
2	2	1.815584e+06
3	3	5.312600e+06
4	4	5.493380e+06

```

from sklearn.model_selection import train_test_split


```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)

regressor_accuracy = LinearRegression()
regressor_accuracy.fit(X_train, y_train)


y_pred = regressor_accuracy.predict(X_test)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R squared score:', metrics.r2_score(y_test, y_pred))
```

 Mean Absolute Error: 3004513.8682692987  
 Mean Squared Error: 19196941411388.652  
 Root Mean Squared Error: 4381431.434062235  
 R squared score: -0.8044963011047319

```
#Implementing Logistic Regression
from sklearn.linear_model import LogisticRegression
reg = LogisticRegression()
reg.fit(X, Y)
```

```
test_predicted_lreg = pd.DataFrame()
test_predicted_lreg['Id'] = data_test.Id
test_predicted_lreg['Prediction'] = reg.predict(data_test.drop('Id', axis=1))
test_predicted_lreg.head()
```

 /usr/local/lib/python3.6/dist-packages/sklearn/linear\_model/\_logistic.py:940: Converge  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

	<b>Id</b>	<b>Prediction</b>
<b>0</b>	0	4888774.0
<b>1</b>	1	4067566.0
<b>2</b>	2	3426169.0
<b>3</b>	3	5017319.0
<b>4</b>	4	3752885.0

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)

regressor_accuracy = LogisticRegression()
regressor_accuracy.fit(X_train, y_train)

y_pred = regressor_accuracy.predict(X_test)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
```

```
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R squared score:', metrics.r2_score(y_test, y_pred))
```



Mean Absolute Error: 2065004.2142857143  
 Mean Squared Error: 13982633811870.072  
 Root Mean Squared Error: 3739336.012164469  
 R squared score: -0.31435578473208037  
 /usr/local/lib/python3.6/dist-packages/sklearn/linear\_model/\_logistic.py:940: Converge  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

#Implementing Support Vector Machines

```
from sklearn.svm import SVC
classifier = SVC(kernel='rbf', random_state = 1)
classifier.fit(X, Y)
```

```
test_predicted_svm = pd.DataFrame()
test_predicted_svm['Id'] = data_test.Id
test_predicted_svm['Prediction'] = classifier.predict(data_test.drop('Id', axis=1))
test_predicted_svm.head()
```



	Id	Prediction
0	0	5461700.0
1	1	2344689.0
2	2	1904842.0
3	3	2371202.0
4	4	1619683.0

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)
```

```
regressor_accuracy = SVC(kernel='rbf', random_state = 1)
regressor_accuracy.fit(X_train, y_train)
```

```
y_pred = regressor_accuracy.predict(X_test)
```

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R squared score:', metrics.r2_score(y_test, y_pred))
```



Mean Absolute Error: 2357957.785714286  
 Mean Squared Error: 9352554362343.428  
 Root Mean Squared Error: 3058194.624667212  
 R squared score: 0.12086777830568829

```
#Implementing Ridge and Lasso Model
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
from sklearn import metrics
```

```
#Lasso Regression
```

```
model = Lasso(alpha=5.5)
```

```
model.fit(X, Y)
```

```
test_predicted = pd.DataFrame()
```

```
test_predicted['Id'] = data_test.Id
```

```
test_predicted['Prediction'] = model.predict(data_test.drop('Id', axis=1))
```

```
test_predicted.head()
```



```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_coordinate_descent.py:47:
positive)
```

	<b>Id</b>	<b>Prediction</b>
<b>0</b>	0	4.669774e+06
<b>1</b>	1	2.740680e+06
<b>2</b>	2	1.815398e+06
<b>3</b>	3	5.311360e+06
<b>4</b>	4	5.492584e+06

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)
```

```
regressor_accuracy = Lasso(alpha=5.5)
```

```
regressor_accuracy.fit(X_train, y_train)
```

```
y_pred = regressor_accuracy.predict(X_test)
```

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
```

```
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
```

```
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
print('R squared score:', metrics.r2_score(y_test, y_pred))
```



```
Mean Absolute Error: 3004980.5481886053
```

```
Mean Squared Error: 19201868342792.414
```

```
Root Mean Squared Error: 4381993.649332734
```

```
R squared score: -0.8049594285010895
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_coordinate_descent.py:47:
positive)
```

```
#Ridge Regression
```

```
model = Ridge(alpha=330)
```

```
model.fit(X, Y)
```

```
test_predicted_ridge = pd.DataFrame()
```

```
test_predicted_ridge['Id'] = data_test.Id
test_predicted_ridge['Prediction'] = model.predict(data_test.drop('Id', axis=1))
test_predicted_ridge.head()
```



	Id	Prediction
0	0	4.229698e+06
1	1	3.822859e+06
2	2	3.636879e+06
3	3	3.882530e+06
4	4	3.957615e+06

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)

regressor_accuracy = Ridge(alpha=330)
regressor_accuracy.fit(X_train, y_train)

y_pred = regressor_accuracy.predict(X_test)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R squared score:', metrics.r2_score(y_test, y_pred))
```



```
Mean Absolute Error: 1930832.797409726
Mean Squared Error: 10032310013899.816
Root Mean Squared Error: 3167382.202055795
R squared score: 0.05697132039594843
```

```
#Random Forest Regression Implementation
from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(n_estimators=150)
model.fit(X, Y)

test_predicted_forest = pd.DataFrame()
test_predicted_forest['Id'] = data_test.Id
test_predicted_forest['Prediction'] = model.predict(data_test.drop('Id', axis=1))
test_predicted_forest.head()
```



	Id	Prediction
0	0	3.991361e+06
1	1	3.421446e+06
2	2	3.575349e+06
3	3	3.433682e+06
4	4	4.420646e+06

```


from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)

regressor_accuracy = RandomForestRegressor(n_estimators=150)
regressor_accuracy.fit(X_train, y_train)

y_pred = regressor_accuracy.predict(X_test)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R squared score:', metrics.r2_score(y_test, y_pred))

```


 Mean Absolute Error: 1563027.243333333  
 Mean Squared Error: 9914139832189.248  
 Root Mean Squared Error: 3148672.7096014996  
 R squared score: 0.06807921780668336

```

#Decision tree Regression Model
from sklearn.tree import DecisionTreeRegressor
regressor = DecisionTreeRegressor(random_state = 0)
regressor.fit(X, Y)

test_predicted_tree = pd.DataFrame()
test_predicted_tree['Id'] = data_test.Id
test_predicted_tree['Prediction'] = model.predict(data_test.drop('Id', axis=1))
test_predicted_tree.head()

```



	Id	Prediction
0	0	3.991361e+06
1	1	3.421446e+06
2	2	3.575349e+06
3	3	3.433682e+06
4	4	4.420646e+06

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=0)

regressor_accuracy = DecisionTreeRegressor(random_state = 0)
regressor_accuracy.fit(X_train, y_train)

y_pred = regressor_accuracy.predict(X_test)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R squared score:', metrics.r2_score(y_test, y_pred))

```



```
Mean Absolute Error: 1917038.7142857143
Mean Squared Error: 10582062293044.715
Root Mean Squared Error: 3253008.1913583796
R-squared: 0.005305060843183531
```